

--	--	--

December 29, 2024

C964: Computer Science Capstone

Movie Recommendation Service

Task 2 parts A, B, C and D

Part A: Letter of Transmittal.....	2
Part B: Project Proposal Plan.....	5
Project Summary.....	5
Data Summary.....	6
Implementation.....	6
Timeline.....	7
Evaluation Plan.....	8
Resources and Costs.....	8
Part C: Application.....	10
Part D: Post-implementation Report.....	11
Solution Summary.....	11
Data Summary.....	11
Machine Learning.....	13
Validation.....	15
Visualizations.....	16
User Guide.....	18

--	--	--

Part A: Letter of Transmittal

December 29, 2014

To Joey Wheeler, Lunabox CEO

834 Philmont Street

Brooklyn, NY 11214

Dear Mr. Wheeler,

I am writing to you today to present a new movie recommendation service for your company Lunabox with the goal of enhancing your user base's experience. We aim to increase user engagement by utilizing advanced algorithms that will be capable of providing personalized movie recommendations for each user that are tailored to their personal preferences. In doing so, we will create a more engaging and captivating viewing experience that will keep users satisfied, thus improving user retention.

In the current state of your product, Lunabox's streaming service lacks the ability to make suggestions for what to watch next. This inability to personalize the experience to each user may be discouraging to users who find difficulty in making use of the streaming service. As a result, users may be more inclined to cancel their memberships if they do not receive much value out of it.

The proposed solution will be a movie recommendation service that gathers user data and feedback on movies. This data will be used to tailor its recommendations to users. When a user searches up a particular movie that they enjoy, a list of movies that have rated positively by users with similar preferences will be returned to the user. This will be done through the use of machine learning algorithms that will continuously learn from the interactions of Lunabox's user

	2	
--	---	--

--	--	--

base over time. This solution aims to increase the usage of the streaming service, thus improving user satisfaction.

Should this solution be implemented, the recommendation service will provide many benefits to the organization. Some of the key benefits are listed below:

1. By offering a tailored user experience, users will log more activity on the streaming service, thus enhancing user engagement.
2. As users will be utilizing Lunabox more, they are more likely to see value in their subscription making them more likely to renew their membership. This in turn increases user retention.
3. Lunabox's direct competitors currently lack the ability to recommend movies to its user base. Implementing this gives them a direct edge over its competition, thus making Lunabox more appealing to prospective users.

Developing this recommendation system will require a software engineer, a machine learning specialist, and a quality assurance engineer on staff. Below is a brief summary of costs, timelines, and ethical concerns regarding the development process.

Costs: The cost of implementing the solution will be approximately \$21,675 for a staff of three people comprising of a software engineer, machine learning specialist, and a quality assurance engineer each earning \$85/hour.

Timelines: The estimated time to complete this project will be 85 hours of work with a product delivery date of 01/31/2025.

Ethical Concerns: Because this recommendation service tracks user data, it is important that we comply with data compliance laws and protect the handling of user data.

	3	
--	---	--

--	--	--

Me and my team have a long history of experience in software development as well as machine learning and are confident and highly motivated in delivering a product that will meet the highest of standards for both Lunabox's team as well as its users. Having led many projects similar to this, I believe that I will efficiently and reliably meet all timelines as well as keep development costs within the allotted budget.

I look forward to talking to you further regarding this proposal! If you need to contact me, my phone number is listed below.

Sincerely,

Luigi Celestino

Project Manager

(123) - 456 - 7890

	4	
--	---	--

--	--	--

Part B: Project Proposal Plan

Project Summary

Lunabox as a streaming service currently has no way of making recommendations to its users for what movie to watch. The only way users can find movies is by already knowing them beforehand and searching it up on the platform. This creates a lackluster experience for users that also fails to differentiate Lunabox from its competitors.

Lunabox, as an organization, is less likely to grow its user base as a result of this. This proposal will outline a solution to resolving this issue, ultimately delivering a recommendation service for users that created tailored experiences that increase user engagement.

The following list will outline the project deliverables:

1. Project plan will be written that outlines each step of the development process.
2. Dataset that will be used.
3. Implementation of the recommendation system using Jupyter Notebook.
4. Interactive user interface.
5. Project documentation that includes a user guide.

This recommendation system will benefit Lunabox in many ways. Initially, the end users will be provided with a superior experience that is tailored to their preferences. As a result of this, user engagement will improve which will improve user retention rates as well. Because of this, users will be more likely to renew their subscriptions which results in higher revenue. Another benefit to Lunabox is that the recommendation system will provide them an edge against their competition, allowing them to stand out.

	5	
--	---	--

--	--	--

Data Summary

The dataset that will be used for this project will be sourced from grouplens.org at the following link: <https://files.grouplens.org/datasets/movielens/ml-latest-small.zip>. This dataset has already been cleaned prior to the developmental process. It will be further cleaned during development as titles will be reformatted to exclude non-alphanumeric characters. The dataset is quite expansive and includes approximately 10,000 films with user-ratings and genres for each film as well as its own unique ID. As such, the movie is well-suited for the needs of the project as it can easily accommodate the movies offered by the Lunabox streaming service. There will be no ethical or legal concerns in the use of this dataset as it is free to use by the public and contains no sensitive information.

Implementation

In the development of this product, we will use the waterfall methodology. Because this project has clearly defined objectives for each step during development and the team already has experience with similar projects, the waterfall methodology with its structured approach is a good methodology to employ.

The project will begin with setting up the development of the project. This will involve the planning of the project as well as setting up the environment that we will develop the product in such as installing the necessary libraries. The dataset will then be acquired and explored to identify patterns and to determine if any modifications to the dataset is required. After the data has been thoroughly explored, we can begin to develop and train the model that will be used

	6	
--	---	--

--	--	--

for the movie recommendation system. Once the model has been developed, a user interface will be implemented that allows end users to input a movie which will then return a list of movie recommendations based on the user input. Finally, testing of the product can take place and the completed product will be made available for submission.

Timeline

Milestone or deliverable	Duration (hours or days)	Projected start date	Anticipated end date
Project plan will be written up	2 days	12/30/2024	01/01/2025
Development environment will be set up	1 day	01/01/2025	01/02/2025
Acquire dataset	1 day	01/02/2025	01/03/2025
Explore the dataset (and modify if necessary)	7 days	01/03/2025	01/10/2025
Development and training of ML model	10 days	01/10/2025	01/20/2025
Implementation of interactive user interface	3 days	01/20/2025	01/23/2025
QA Testing	4 days	01/23/2025	01/27/2025
Final Fixes	3 days	01/27/2025	01/30/2025
Product Submission	1 day	01/30/2025	01/31/2025

	7	
--	---	--

--	--	--

Evaluation Plan

During the development of this product, proper version control practices will be used to ensure that no faulty code makes its way to the live product. This will be done via GitHub to make sure that each chunk of code written by the team can be tested in its own independent environment that each team member will be working in.

Validation for the performance of the product will be evaluated in the future by users of Lunabox's streaming service once the product goes live. Because the product is supposed to recommend movies that the user will like, it is only possible to receive quantifiable data regarding the success of this directly from Lunabox's user base. This will be done by prompting users to rate the quality of the recommendations given to them. This user feedback will be logged and stored by the service. If less than 50% of users are satisfied with their automated recommendations, the service will be deemed unsuccessful and adjustments to the model will be required.

Resources and Costs

The product will be developed entirely on Jupyter Notebook using Python 3. Because Jupyter Notebook is free, the only hardware/software costs come from the work stations for the three team members. Each work station will cost \$1,500 each, or a total of \$4,500. Each team member on staff earns an hourly rate of \$85/hour. With an expected project duration of a total of 255 work hours, the cost of labor will cost approximately \$21,675. As a result of being hosted entirely on the Jupyter Notebook as a standalone product, it will cost nothing to be deployed. Overall, the total cost of the entire project amounts to \$26,175.

	8	
--	---	--

--	--	--

Resource	Description	Cost
Software Engineer	Will be responsible for the front and backend	\$7,225
Machine Learning Specialist	Handles the development of the ML model	\$7,225
QA Engineer	Ensures the product passes all required testing	\$7,225
Work stations	Three work stations, each costing \$1,500	\$4,500
Total	Cost of the entire project	\$26,175

	9	
--	---	--

--	--	--

Part C: Application

The application will be included as part of the submission in the folder.

*Because the application is run locally from a Jupyter Notebook, there are no security concerns and as such, no industry appropriate security features were implemented.

	10	
--	----	--

--	--	--

Part D: Post-implementation Report

Create a post-implementation as outlined below. Provide sufficient detail so that a reader knowledgeable in computer science but unfamiliar with your project can understand what you have accomplished. Using examples and visualizations (including screenshots) beyond the three required is recommended (but not required). **Write everything in the past tense.**

Solution Summary

The problem that Lunabox faced was the lack of a movie recommendation system for its users. As a result, users were faced with a lackluster experience. Without this feature, users were required to have their own knowledge of which movies to watch to discover new content. As a result, user engagement was not as high compared to their competitors. As a result, Lunabox contacted us to fix this issue. The solution that was created was a standalone tool that used machine learning algorithms to recommend to its users different movies to watch based on their preferences. It provided users with an interactive user interface that allowed for users to input a movie that they liked and through the use of user-based collaborative filtering, would output a list of movies that it recommended. The application provided a solution to Lunabox's issue by providing a feature that recommended movies to Lunabox's users.

Data Summary

The dataset used for this project were CSV files describing movies and their features such as genres, as well as their user-reported ratings. It was sourced from grouplens' website, a research lab specializing in recommender systems. The data had already been cleaned prior to

	11	
--	----	--

--	--	--

development as it was handled by grouplens' researchers. Nonetheless, exploration of the data still took place to determine if there were any anomalies or missing values. It was determined that the data did not have any null values, nor were there any anomalies that would require further investigation. However, further data cleaning did take place during development as the titles of the movies needed to be reformatted to provide more reliable search queries to the user. Figure 1 shows a section of the movies dataset that was used in the project while figure 2 shows the ratings dataset that corresponds to the movie dataset on their movied.

movied		title		genres	format_title
0	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy		Toy Story 1995
1	2	Jumanji (1995)	Adventure Children Fantasy		Jumanji 1995
2	3	Grumpier Old Men (1995)	Comedy Romance		Grumpier Old Men 1995
3	4	Waiting to Exhale (1995)	Comedy Drama Romance		Waiting to Exhale 1995
4	5	Father of the Bride Part II (1995)	Comedy		Father of the Bride Part II 1995
...
9737	193581	Black Butler: Book of the Atlantic (2017)	Action Animation Comedy Fantasy		Black Butler Book of the Atlantic 2017
9738	193583	No Game No Life: Zero (2017)	Animation Comedy Fantasy		No Game No Life Zero 2017
9739	193585	Flint (2017)	Drama		Flint 2017
9740	193587	Bungo Stray Dogs: Dead Apple (2018)	Action Animation		Bungo Stray Dogs Dead Apple 2018
9741	193609	Andrew Dice Clay: Dice Rules (1991)	Comedy		Andrew Dice Clay Dice Rules 1991
9742 rows x 4 columns					

Figure 1: The movies dataset

--	--	--

	userId	movieId	rating	timestamp
0	1	1	4.0	964982703
1	1	3	4.0	964981247
2	1	6	4.0	964982224
3	1	47	5.0	964983815
4	1	50	5.0	964982931
...
100831	610	166534	4.0	1493848402
100832	610	168248	5.0	1493850091
100833	610	168250	5.0	1494273047
100834	610	168252	5.0	1493846352
100835	610	170875	3.0	1493846415
100836 rows × 4 columns				

Figure 2: The ratings dataset

Machine Learning

The recommendation system that was developed used user-based collaborative filtering and utilized cosine similarity scores to compute the similarity between the term that the user entered and the movies in the database. In order for this to happen, the movie titles were first vectorized so that cosine similarity could be used. Cosine similarity is a metric used to determine how similar two vectors are, or in this case two movies. This calculation was done through the usage of the sklearn library and returns a search result of the most similar movies.

```

from sklearn.metrics.pairwise import cosine_similarity
import numpy as np

# Cosine similarity will allow us to compute similarity between terms that are entered
def search(title):
    title = format_title(title)
    query_vec = vectorizer.transform([title])
    similarity = cosine_similarity(query_vec, tfidf).flatten()
    indices = np.argpartition(similarity, -5)[-5:]
    results = movies.iloc[indices][::-1]
    return results

```

Figure 3: Search function that uses cosine similarity to return similar movies

Upon returning a list of similar movies, the program determined a list of users from the database similar to the user of the tool. This was calculated by matching the movie ID that the user input to other unique users who rated that movie at least four stars.

```

similar_users = ratings[(ratings["movieId"] == movie_id) & (ratings["rating"] >= 4)]["userId"].unique()

similar_users

array([ 1,  5,  7, 17, 19, 31, 40, 43, 45, 46, 57, 63, 64,
        66, 71, 73, 78, 86, 91, 96, 98, 103, 107, 121, 124, 135,
       137, 141, 145, 151, 156, 159, 160, 161, 166, 169, 171, 177, 178,
       179, 182, 185, 186, 191, 201, 202, 206, 217, 220, 229, 234, 239,
       240, 247, 249, 252, 254, 263, 264, 269, 270, 273, 274, 275, 276,
       277, 280, 282, 288, 290, 291, 292, 304, 307, 328, 330, 332, 336,
       337, 339, 341, 347, 350, 353, 357, 359, 364, 367, 378, 380, 382,
       385, 389, 396, 399, 411, 414, 420, 422, 434, 436, 438, 443, 448,
       451, 453, 456, 460, 468, 469, 470, 471, 474, 476, 477, 483, 484,
       488, 492, 500, 504, 509, 514, 517, 524, 525, 533, 534, 550, 555,
       559, 561, 562, 570, 572, 573, 579, 584, 587, 590, 596, 597, 601,
       603, 605, 607, 610], dtype=int64)

similar_user_recs = ratings[(ratings["userId"].isin(similar_users)) & (ratings["rating"] >= 4)]["movieId"]

similar_user_recs

0      1
1      3
2      6
3     47
4     50
...
100830 166528
100831 166534
100832 168248
100833 168250
100834 168252
Name: movieId, Length: 20964, dtype: int64

```

Figure 4: Returning the list of similar users based on the movie input

--	--	--

The list was then further refined to only include movie recommendations that had at least 20% of users who were deemed similar to the user of the tool. Furthermore, the similarity score was converted to a percentage to easily represent how strong of a recommendation the movie was.

```
# Converts the similarity count to a percentage of similar users that recommend that movie
similar_user_recs = similar_user_recs.value_counts() / len(similar_users)

# Returns movie recommendations that had more than 20% of similar users who recommended it
similar_user_recs = similar_user_recs[similar_user_recs > .2]
```

Figure 5: Similarity score was converted to a percentage and the list of movies was further refined to only include movies with at least 20% of four star reviews from similar users

Following this, the program would then filter the results even further until the recommendation tool ultimately outputs a list of movies that it recommends the user watch based on their input.

Cosine similarity was a good fit for the implementation of the system because it would recommend movies that had similar names to the movie that the user was interested in. This is an obvious choice because naturally, if a person enjoyed a movie from a franchise, a sequel which bears the same franchise name should be recommended to the user. Collaborative filtering was also a good fit for this tool because we wanted the user's input to be tailored to their preferences so identifying patterns in movie preferences from similar users helped to create the recommendation tool.

Validation

Because the performance metric of the recommendation system is gauged by user satisfaction, the means of validating the success of the tool is to directly request for feedback from the user. This will be done post-implementation when the product is already live for Lunabox's users to utilize. Following the usage of the tool, the user will be prompted for feedback regarding the

	15	
--	----	--

--	--	--

quality of the recommendations and their feedback will be stored for analysis. If the satisfaction rate reported by the users who used the tool are less than the targeted 50%, the tool will be considered a failure and the system will need to be readjusted.

Visualizations

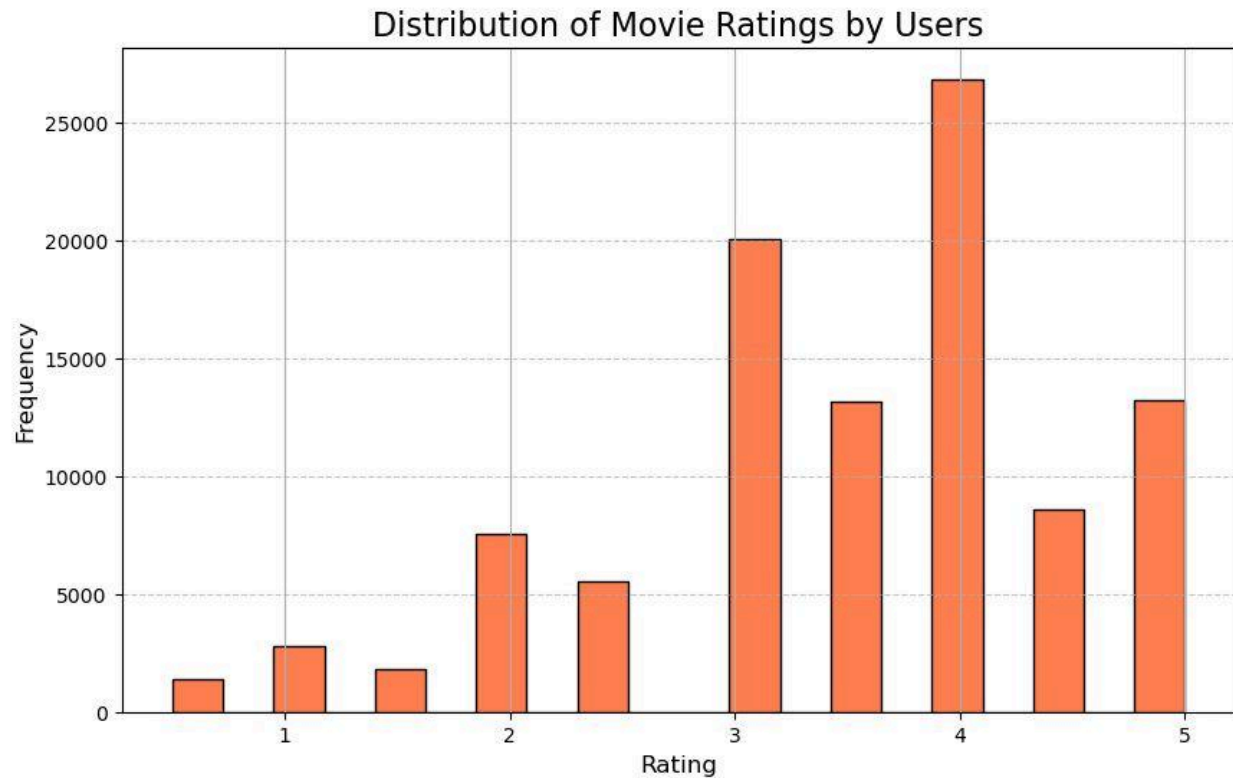


Figure 6: Bar chart of distribution of movie ratings reported by users

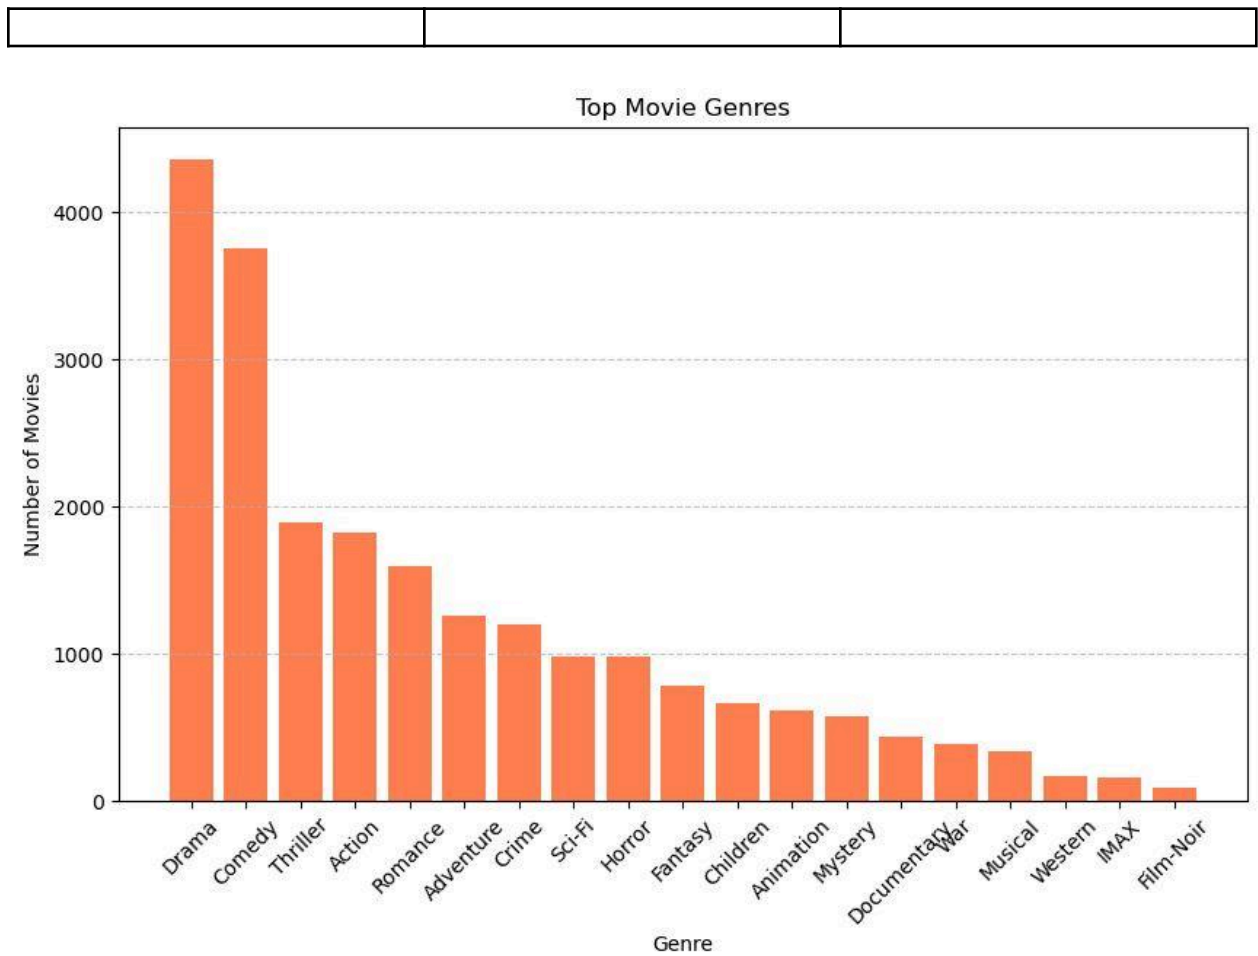


Figure 7: Bar chart of the top genres from the dataset

--	--	--

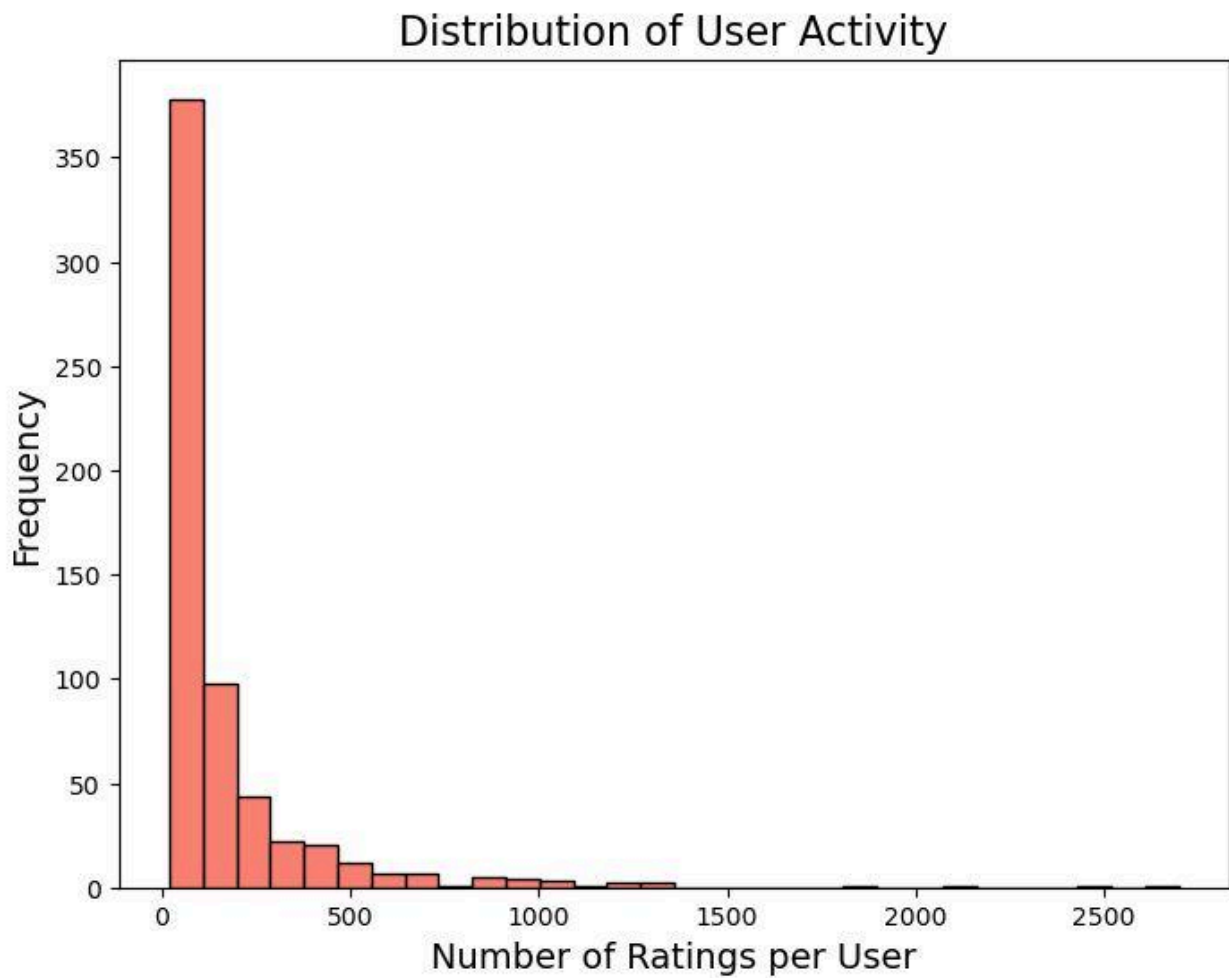


Figure 8: Histogram showcasing the distribution of user activity based on their number of ratings

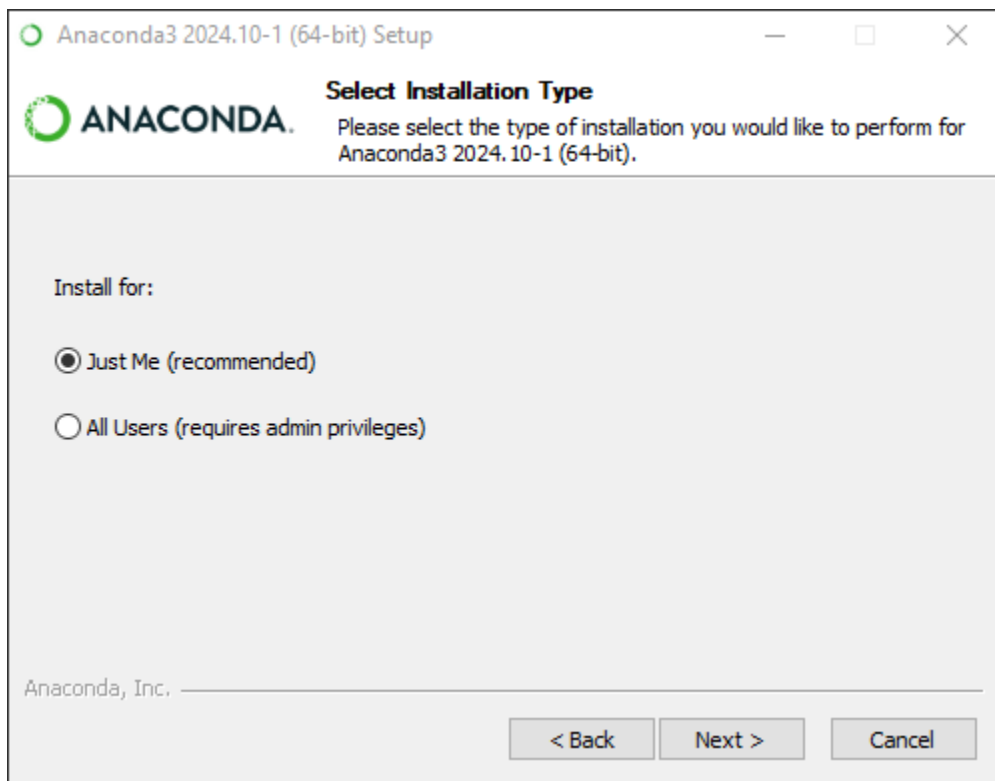
User Guide

1. Download the folder named C964 Capstone Submission.
 - a. Inside are the datasets for the movies and the ratings as well as the Jupyter Notebook file for the standalone tool.

--	--	--

2. Navigate to <https://www.anaconda.com/download/success> to download Anaconda if you don't have it yet.

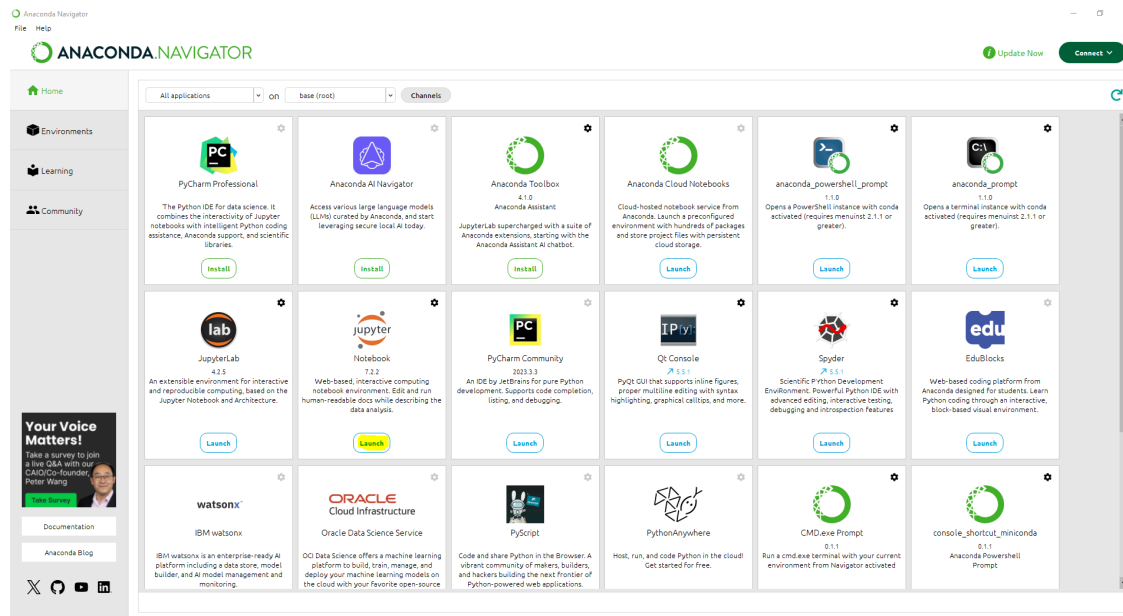
3. Run the installer and when prompted to select your installation type, select "Just Me (recommended)"



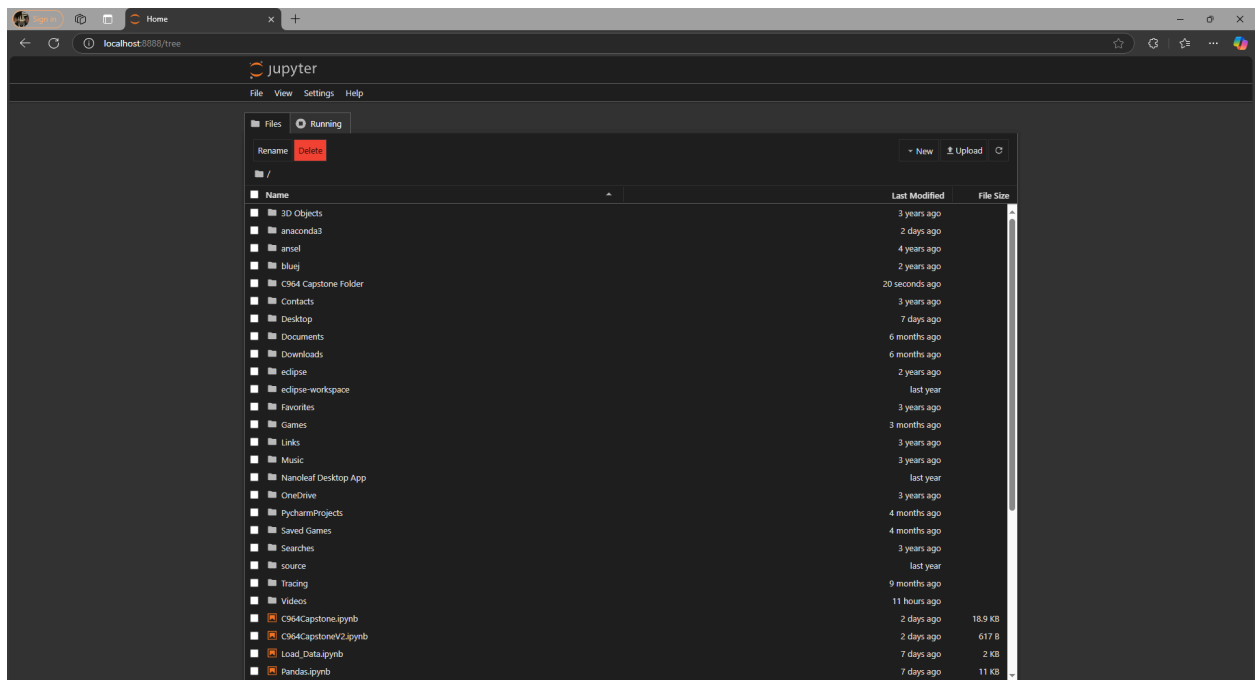
4. Click through the prompts and let the installer run.

	19	
--	----	--

5. Once Anaconda Navigator is installed, make sure that Notebook is installed from the home screen. Once installed, click launch.



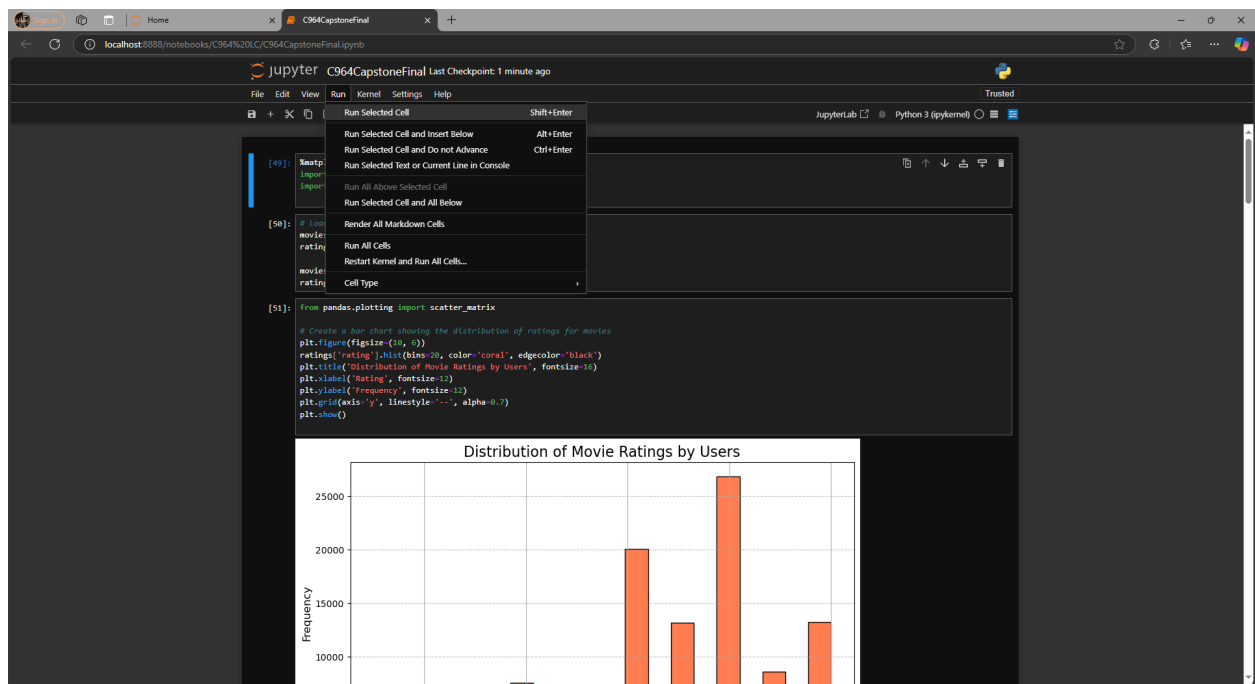
6. Your internet browser will open Jupyter Notebook and your screen will appear like this.



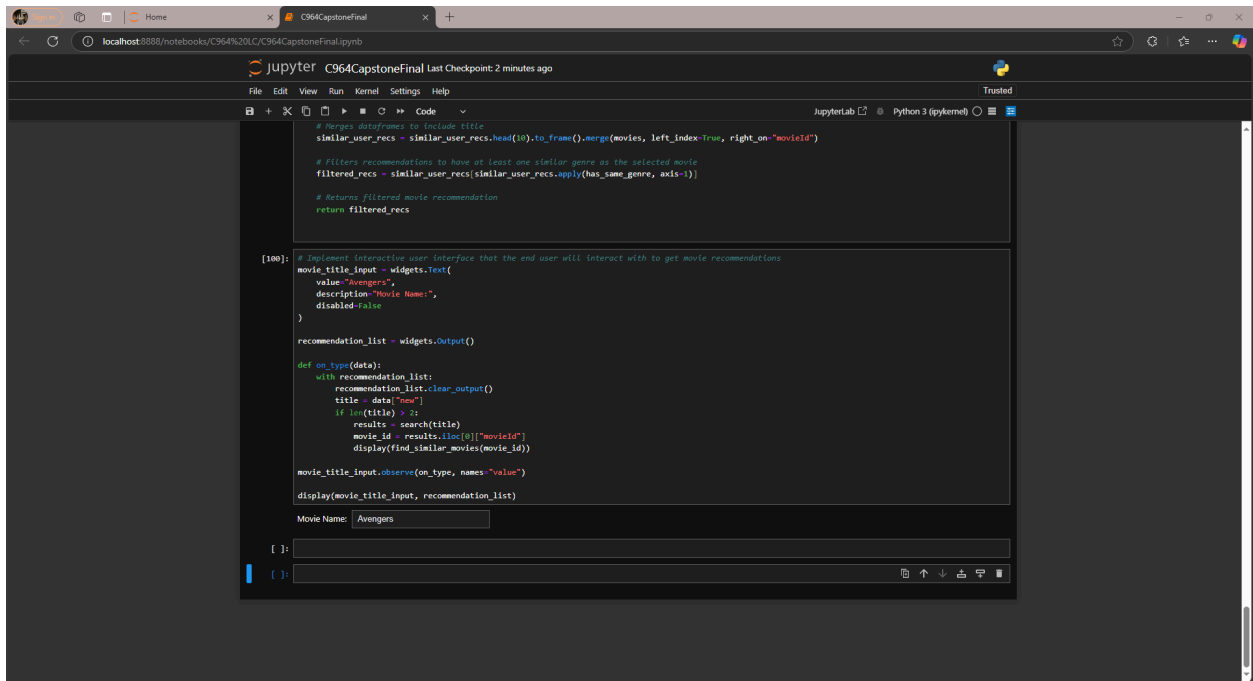
--	--	--

7. Once here, navigate through your file directory to where the folder was downloaded to and open the folder. Double click the C964CapstoneFinal.ipynb file.

8. At the top, click Run -> Run All Cells. Because Anaconda downloads packages for you, no additional libraries need to be manually downloaded.



9. The Notebook will redirect you to the bottom where the user will interact with.



The screenshot shows a JupyterLab notebook titled "C964CapstoneFinal" with a last checkpoint of 2 minutes ago. The interface includes a top bar with navigation icons and a menu bar with options like File, Edit, View, Run, Kernel, Settings, and Help. The main area displays a code editor with the following Python code:

```
# Merges dataframes to include title
similar_user_recs = similar_user_recs.head(10).to_frame().merge(movies, left_index=True, right_on="movieId")

# Filters recommendations by those at least one similar genre as the selected movie
filtered_recs = similar_user_recs[similar_user_recs.apply(lambda row: row['genre'] in row['similar_genres'], axis=1)]

# Returns filtered movie recommendation
return filtered_recs
```

Below the code editor, there is an interactive user interface. It starts with a text input field labeled "Movie Name:" containing the text "Avengers". Below this is a button labeled "Recommend". To the right of the input field is a scrollable list of movie titles, with "Avengers: Endgame" selected. At the bottom, there is a section labeled "Recommendations" which is currently empty.

--	--	--

10. This bottom search bar labeled “Movie Name:” is where the user will input the name of a movie and the tool will output a list of movie recommendations.

The screenshot shows a JupyterLab environment with a Python 3 kernel. The code in the cell implements an interactive user interface for movie recommendations. It uses a text input widget for the movie name and an output widget for the recommendations. The function `on_type` is triggered when the input value changes, and it searches for similar movies based on the input title.

```
[100]: # Implement interactive user interface that the end user will interact with to get movie recommendations
movie_title_input = widgets.Text(
    value="Avengers",
    description="Movie Name:",
    disabled=False
)

recommendation_list = widgets.Output()

def on_type(data):
    with recommendation_list:
        recommendation_list.clear_output()
        title = data["new"]
        if len(title) > 2:
            results = search(title)
            movie_id = results.iloc[0]["movieid"]
            display(find_similar_movies(movie_id))

movie_title_input.observe(on_type, names="value")
display(movie_title_input, recommendation_list)
```

Movie Name:

	count	movieid	title	genres	format_title
990	1.000000	1291	Indiana Jones and the Last Crusade (1989)	Action Adventure	Indiana Jones and the Last Crusade 1989
900	0.884298	1198	Raiders of the Lost Ark (Indiana Jones and the...)	Action Adventure	Raiders of the Lost Ark Indiana Jones and the ...
224	0.801653	260	Star Wars: Episode IV - A New Hope (1977)	Action Adventure Sci-Fi	Star Wars Episode IV A New Hope 1977
898	0.793388	1196	Star Wars: Episode V - The Empire Strikes Back...	Action Adventure Sci-Fi	Star Wars Episode V The Empire Strikes Back 1980
911	0.719008	1210	Star Wars: Episode VI - Return of the Jedi (1983)	Action Adventure Sci-Fi	Star Wars Episode VI Return of the Jedi 1983
969	0.661157	1270	Back to the Future (1985)	Adventure Comedy Sci-Fi	Back to the Future 1985
257	0.628099	296	Pulp Fiction (1994)	Comedy Crime Drama Thriller	Pulp Fiction 1994
314	0.603306	356	Forrest Gump (1994)	Comedy Drama Romance War	Forrest Gump 1994

[]:

[]: